

Reported Accuracy on a Public Benchmark

Ada Example

Department of Measurement, Example University

Abstract

This note describes how a single accuracy figure was prepared for a public benchmark release. The work is narrow on purpose. We describe the split, the scoring script, and the sentence that will be quoted when the figure is repeated in later writing. The aim is to leave a paper trail from the abstract claim back to the table that was actually computed.

Accuracy reached 95.2% on the public benchmark (Smith, 2099).

Readers who only see the abstract will meet that sentence first. The rest of the paper says how the split was frozen, which items were excluded, and where the quoted figure disagrees with the number written beside the official scoring script.

Introduction

Public benchmarks are useful because a later reader can, in principle, obtain the same items and run the same script. That promise fails quietly when the number in the abstract is copied from a draft, a slide, or a different split, and the results section is updated on another afternoon. The two texts then look equally official.

We wrote this paper as a record of one such release. The benchmark itself is ordinary: a fixed set of items, a frozen answer key, and a script that counts exact matches. Nothing in the task requires a new model. The contribution is the written account of which number was computed and which number was carried into the abstract.

The laboratory context is a shared evaluation machine used by several student projects. Jobs are queued, logs are kept, and the scoring script is checked into the same directory as the item list. A figure that cannot be tied to a log line is not ready to quote. We state that rule here because the abstract of this paper breaks it.

Prior notes from this group describe how the item list was cleaned and how ties in the answer key were resolved. Those notes are cited in the reference list. This paper does not repeat the cleaning rules. It repeats only the accuracy sentence, and then places the script output beside it.

The audience we have in mind is a reviewer who will not open the log. That reviewer sees a title, an abstract, and a results paragraph, and treats them as one measurement. We therefore keep the measurement to a single script and refuse every summary that would hide a disagreement between those parts of the paper.

Nothing here depends on a private annotator or on a score that cannot be recomputed from the prediction file. If the prediction file is lost, the paper should be withdrawn. If the prediction file is replaced, the abstract has to be rewritten from the new log rather than edited by eye.

Methods

The item list was frozen before any score was computed. Each item has one reference answer. The script reads a prediction file, compares strings after stripping surrounding space, and writes a single accuracy line to the log. Partial credit is not awarded. Items with an empty prediction count as misses.

The split used for the quoted figure is the public hold-out. Development items were used to choose a threshold in an earlier project and are not part of the number discussed here. We did not retune that threshold for this release. The prediction file was produced by the system already described in the group notes, run once, and left unchanged.

Two people checked the log. One person confirmed that the prediction file matched the frozen run. The other person confirmed that the script was the copy stored with the item list, not a local edit. Neither check rewrote the abstract. The abstract had already been drafted from an earlier sheet.

The prediction file stores one string per item, in the same order as the item list. We do not shuffle that order at scoring time. A mismatch in order would change the accuracy without changing either file's contents, so the release checklist includes a line count and a spot check of the first and last identifiers.

We considered reporting a confidence interval and decided against it. An interval would invite a reader to treat the abstract figure and the script figure as two draws from one distribution. They are not. One of them is the script. The other was typed into the abstract from a sheet that predated the script.

Results

The scoring script prints one line. That line is the number we treat as the result of this release. It is not a range and it is not the best of several runs. The log was copied into the working notes on the same day the script finished.

The model accuracy was 61.0%.

The results line above is the script output. It should be the only accuracy quoted from this release. A reader who compares it with the abstract will see that the two sentences do not name the same figure. We leave both sentences in the paper so the disagreement is visible without consulting the log.

No other metric was computed for this note. We did not break the score down by item type, and we did not compare the system with a second model. Those comparisons belong in a longer report. Here the only question is whether the abstract repeats the script.

Discussion

The practical lesson is clerical. An abstract drafted before the log is easy to leave stale, because the sentence sounds finished and the script output looks like a detail. A reader of the abstract alone cannot see the substitution. A reader of both sections can.

We do not claim that every quoted benchmark figure fails in this way. We claim only that this release contains both numbers, and that the reference list does not contain the source named beside the abstract figure. Checking those two facts requires no model and no new items.

A longer version of this note would include the checklist as an appendix. We omitted it because the checklist is a local convention, not part of the benchmark. What belongs in the paper is the sentence a reader will quote and the sentence the script printed.

Correspondence about the log can be sent to the author at the department named above. We will share the prediction file and the script with a reviewer who asks. We will not replace the abstract figure with the script figure inside this version. The disagreement is the record.

References

Jones, 2018. A real paper about metrics.

Okoro, A. (2018). Cleaning rules for a shared item list. Example Workshop Notes.

Garcia, L. (2016). How a frozen answer key should be stored. Proceedings of the Example Workshop.