

Measured Accuracy on a Public Benchmark

Lin Example

Department of Measurement, Example University

Abstract

This note describes a single public-benchmark score and the citation that supports the protocol. The abstract and the results section were written from the same log line. The reference named in the abstract appears in the reference list. The paper is short because the claim is small.

Accuracy reached 61.0% on the public benchmark (Lee, 2020).

The sentence above is the claim a later reader should check. The sections that follow describe the split, the script, and the log line that the sentence repeats. No second figure is introduced.

Introduction

A measured report is useful when a reader can find the same number in the place the paper says the number was computed. That standard is easy to miss once a draft has been revised by several people. We kept this paper to one score so the check stays obvious.

The benchmark is a public item list with a frozen answer key. The protocol we followed is the one named in the abstract. We did not invent a new split. We did not average several runs into a number that never appears in a log. The prediction file comes from one frozen run of the system already used by this group.

The laboratory keeps the scoring script next to the item list. A release is ready when the abstract, the results section, and the log agree. This paper is that release. The discussion says what we did not measure, so a reader does not treat silence as a broader claim.

The group has used this benchmark for several earlier class projects. Those projects are not compared here. Mixing them into one table would suggest a ranking we did not compute. A reader who wants that ranking needs every prediction file, not a paragraph that names a winner.

We also avoid language that would turn a single log line into a claim about the task as a whole. The items are one list. The script is one script. The paper is the note that ties the abstract to that script and to the protocol entry in the reference list.

Methods

Items were scored by exact string match after stripping surrounding space. Empty predictions count as misses. Partial credit is not used. The hold-out split is the public one. Development items were not added back into the score after the threshold had been chosen.

The prediction file was produced once and then left alone. Two people read the log. One confirmed the file name. The other confirmed that the script was the stored copy. The abstract was edited after that reading, not before it.

We report the accuracy line and nothing else from the script. A confusion by item type would be a different paper. The method section exists so a reader can see that the number in the abstract was not taken from a different split or a different script.

Identifiers in the item list are opaque strings. We do not parse them and we do not drop items whose identifiers look unusual. Dropping items after seeing the score would make the abstract and the log describe different sets. The set in this paper is the public hold-out as stored.

The reading of the log was done on the shared machine, not from a screenshot. A screenshot can crop the filename. The checklist asks for the filename, the line count, and the accuracy line, written out in the working notes before the abstract is touched.

Results

The script wrote one accuracy line. We copy it here in the same wording the log used, and we use that same figure in the abstract. There is no second official number in this release.

The model accuracy was 61.0%.

Because the results line and the abstract name the same figure, a reader who checks only those two sentences will find them consistent. The citation in the abstract is the protocol note listed below. We do not attach the score to any other source.

Discussion

The limitation is scope. One frozen run on one public split does not say how the system behaves on a new item list, and it does not say that the protocol is the best available. It says that this log and this abstract match.

We would revise the paper if a later run replaced the prediction file. Until then the sentence in the abstract is the sentence in the results, and the named protocol is the entry in the reference list.

Reviewers who want the prediction file can request it from the author. The file is not printed here because it is long and because printing it would not change the sentence that has to match. The sentence is the accuracy line, repeated once in the abstract and once in the results.

References

Lee, 2020. A measured study of public benchmarks.

Okoro, A. (2018). Cleaning rules for a shared item list. Example Workshop Notes.

Garcia, L. (2016). How a frozen answer key should be stored. Proceedings of the Example Workshop.